

Simplifying Python Dependency Management with uv

Managing dependencies in Python can be one of the trickiest parts for beginners learning to write and run scripts. Without the right approach, you might run into “**dependency hell**” – scenarios with conflicting package versions, missing modules, or broken environments. In this comprehensive guide, we’ll explore the challenges of Python dependency management and show how the new `uv` tool helps solve many of these issues. We’ll cover how to set up your environment on both macOS and Windows (even if you don’t have Python installed yet), how to use `uv` in contexts like Jupyter Notebooks and VS Code (or the Cursor editor), and walk through a beginner-friendly tutorial of using `uv` to run standalone Python scripts.

The Challenges of Python Dependency Management

Python’s strength is its rich ecosystem of libraries, but that also means you must manage those library **dependencies** carefully. Some common challenges include:

- **Global vs. Local Environments:** Installing packages “globally” (into your system’s default Python) can lead to version conflicts between projects. For example, one script might need `requests` v2.28 while another needs v2.31 – a single installation can’t satisfy both. The recommended practice is to use isolated **virtual environments** for each project ¹, so each project or script has its own set of packages. However, manually creating and activating these environments can be confusing for newcomers.
- **Multiple Tools and Workflows:** There are many tools to manage Python packages and environments, each with its own workflow:
 - **`pip` + `Virtualenv`:** The built-in solution is to create a virtual environment (using `python -m venv` or `virtualenv`), activate it, then use **`pip`** to install packages. This works but requires several manual steps (creating the env, activating it, remembering to deactivate), and `pip` by itself doesn’t record exact versions unless you manually use `pip freeze` to make a requirements file.
 - **`pipx`:** A tool to install and run Python applications in isolated environments. It’s great for command-line tools, as it creates a virtual env for each tool. However, it’s not designed for managing project-specific libraries; it’s more for running standalone CLI apps.
 - **`Poetry`:** A popular all-in-one package manager that uses a `pyproject.toml` file to declare dependencies and a lockfile for reproducibility. Poetry automatically creates `virtualenvs` and has an easy workflow for adding/removing packages. It improves on traditional tools by handling dependency resolution and locks, but it can be slower and imposes a specific project structure (which some describe as an “opinionated” approach ²).
 - **`Conda`:** An environment manager and package manager that can handle Python packages and even non-Python dependencies (C libraries, etc.). Conda is widely used in data science because it can install complex packages (like scientific libraries with C/Fortran code) from binary channels. It provides **excellent cross-platform consistency** and can manage different Python versions in separate environments ³. However, `conda` operates in its own ecosystem (packages from

conda-forge or Anaconda repositories) and is a relatively heavy installation (the Anaconda distribution is large). Mixing conda and pip can be confusing for beginners, and using conda just for pure-Python projects might be overkill for some cases.

- *(There are other tools like *Pipenv*, *pip-tools*, etc., but the above are the main ones beginners encounter.)*

- **Dependency Resolution and Conflicts:** Tools like pip will install whatever latest version fits the requirement specifiers, but if two packages have incompatible requirements, you get conflicts that are not always straightforward to resolve. Poetry and Conda use more advanced dependency solvers to find a set of packages that work together, but pip's resolver is more basic ⁴. Without a lock file, two installations at different times might resolve to different versions, possibly introducing bugs.

- **Reproducibility:** Ensuring that your code runs the same way on another machine or later in time requires pinning exact package versions. Traditional `requirements.txt` (from pip) is just a list of packages (which may allow newer releases). Tools like *pip-compile* (from *pip-tools*) or Poetry create a **lock file** with exact versions and hashes for reproducible installs. Beginners often skip this, leading to "it worked yesterday, why not today?" issues when a library update introduces changes.

- **Installing Python itself:** A very fundamental challenge – if Python isn't installed on the system, nothing works! On Windows, beginners might not have Python by default. On macOS, Python 2 was historically included but modern Python 3 is not. Setting up Python (and the right version) is the first hurdle. In the past you'd download an installer from *python.org* or use a manager like *pyenv* or Conda to get Python. This adds another layer of complexity, especially if you need a specific Python version for a script.

All these factors contribute to a steep learning curve. It's common for new Python learners to spend hours figuring out environment issues (e.g., "module not found" errors, PATH issues, which Python interpreter VS Code is using, etc.) instead of actually writing code. As one source puts it, the Python ecosystem has "*common pain points*" like slow installs, dependency conflicts, and environment management complexity ⁵.

Traditional Solutions vs. Modern Approach with uv

`uv` is a new-generation Python package and project manager that aims to streamline all these aspects. It is written in Rust and designed as a fast, unified tool to replace or integrate with the traditional tools (pip, virtualenv, pip-tools, pyenv, pipx, etc.) ⁶. To understand how `uv` helps, let's briefly compare it to the traditional approaches:

- **Speed:** `uv` is extremely fast at resolving and installing packages, thanks to its Rust implementation and modern dependency resolver (it uses the PubGrub algorithm under the hood). In fact, it can be *10-100× faster* than pip for certain operations ⁷ ⁵. For example, `uv` can resolve a requirements file of dozens of packages in milliseconds and install them in a fraction of a second ⁸ ⁹. This speed means less waiting and more doing, which is especially nice when you're repeatedly running scripts or updating packages.

- **Integrated Environment Management:** Unlike pip alone, `uv` has **built-in environment management**. It automatically creates and manages virtual environments for you, so you don't

need to manually run `python -m venv` or activate/deactivate environments ¹⁰ ¹¹. Each project or standalone script can run in an isolated environment that `uv` maintains. This isolation prevents the “bleed over” of packages between different projects. As the Real Python introduction notes, “*uv automatically handles virtual environments, creating and managing them as needed to ensure clean and isolated project dependencies.*” ¹⁰ For beginners, this removes a huge source of confusion – you can let `uv` worry about where packages go.

- **Python Version Management:** `uv` can even install and manage multiple Python **versions** on your system, similar to tools like `pyenv`. With a simple command, you can download and set up, say, Python 3.10, 3.11, and 3.12 side by side ¹². You can tell `uv` to run a script with a specific Python version (even PyPy) and it will download it on demand if not already available ¹³ ¹⁴. This means even if you **don’t have Python installed**, `uv` can fetch it for you automatically when needed. According to the documentation, “*Python does not need to be explicitly installed to use uv. By default, uv will automatically download Python versions when they are required.*” ¹⁵. For example, the first time you create an environment or run a script with `uv`, it will pull down the latest Python runtime for you if your system has none. This feature helps newcomers skip the separate Python installation step.

- **Unified Interface:** `uv` combines functionalities that traditionally required separate tools:

- It can act like **pip** (and `pip-tools`): you can use `uv pip install`, `uv pip uninstall`, `uv pip freeze`, or even `uv pip compile` (which is analogous to `pip-tools`’ `compile` to lock dependencies) ⁶ ¹⁶. The benefit is you get those 10–100× speed improvements and some advanced features (like platform-independent lock files, better resolution strategies ⁷).
- It functions like **virtualenv/venv**: creating virtual environments with `uv venv` is simple and lets you easily specify which Python version to use for the env. For instance, `uv venv --python 3.12` will create a new `.venv` using Python 3.12 (downloading 3.12 if necessary) ¹³.
- It covers **pipx** use-cases: `uv` has a concept of “tools”. You can install a command-line tool globally with `uv tool install <pkg>` (similar to `pipx install`) or run a one-off ephemeral command with `uv tool run` (there’s even a shortcut alias `uvx`). For example, `uvx pycowsay "hello world!"` would download and run the `pycowsay` tool in isolation ¹⁷ ¹⁸. This is great for trying out a Python CLI without polluting your main environment.
- It plays the role of **Poetry** (project manager): `uv` can initialize new projects with a `pyproject.toml`, add dependencies (and dev-dependencies) to the project, lock them for reproducibility, and even build and publish packages to PyPI ¹⁹ ²⁰. Essentially, `uv` embraces the modern `pyproject.toml` standard similar to Poetry, but with a focus on performance and flexibility. It supports workspaces, editable installs, version constraints, and more, all via one tool.
- In summary, `uv` is an **all-in-one** solution: “*uv is a Python package and project manager that integrates multiple functionalities into one tool, offering a comprehensive solution for managing Python projects.*” ²¹ It addresses common frustrations by being fast, handling envs automatically, and working cross-platform in a consistent way.
- **Addressing Pain Points:** Because of these capabilities, `uv` directly tackles the earlier challenges:

- **Dependency conflicts and resolution** – `uv` uses a modern solver (PubGrub), so it can quickly find compatible versions or inform you of conflicts with clear error messages. Its resolver is more robust than pip's basic approach, akin to Poetry's solver ⁴.
- **Environment “activation”** – `uv` runs commands in the correct environment automatically. You don't need to activate a virtualenv; `uv run` will ensure the script uses its isolated interpreter and packages. This reduces user error (like installing a package in one env but running in another by accident).
- **Caching and Reusability** – `uv` aggressively caches downloaded packages and built artifacts, so you rarely download the same thing twice ²². The first time you install a package, it's cached (respecting PyPI's HTTP caching headers and other strategies ²³). If you run a script again and nothing changed, `uv` won't reinstall anything – it just reuses the environment. In other words, `uv` will create the env on the first run and **cache it** for subsequent runs of the same script or project. This means you get the convenience of on-demand env creation without the performance penalty on each execution. (If needed, you can explicitly refresh or rebuild, but by default caching makes re-running fast ²⁴.)
- **Works with existing ecosystem** – `uv` doesn't lock you into a proprietary ecosystem. It works with standard Python packages (from PyPI or other indexes), produces standard lockfiles (which are basically enhanced requirements files), and can coexist with tools like pip. You can, for instance, still use pip inside a `uv` environment if you really want, or have `uv` generate a requirements.txt. It's designed to be compatible and flexible ², unlike Conda which has its own package format.

Overall, `uv` “addresses common pain points in the Python ecosystem such as slow installation times, dependency conflicts, and environment management complexity” ⁵ by providing a unified, high-performance tool. Now, let's see how to set it up and use it in practice.

Setting Up Python and uv (Mac & Windows)

Setting up `uv` is straightforward and you can use it even on a system with no Python installed. The process differs slightly between macOS and Windows:

1. **Install `uv` on macOS/Linux:** You can use Homebrew or the provided installation script: - **Homebrew (Mac):** If you have Homebrew, simply run:

```
brew install uv
```

This will fetch the latest `uv` release from Homebrew's core repository ²⁵. - **Installer Script (Mac/Linux):** If you don't have Homebrew, you can use cURL to run the installer. Open a terminal and run:

```
curl -Lsf https://astral.sh/uv/install.sh | sh
```

This downloads and executes the official install script ²⁶. It will detect your OS, download the appropriate `uv` binary, and set it up (it may ask for confirmation to modify your shell profile to add `uv` to PATH). If you don't have curl, the docs note you can use `wget` similarly ²⁷. (The installer is convenient, but as with any script, you can inspect it first by running the URL through `less` or opening it in a browser ²⁸.) - **Alternative:** You can also download the binary from the [GitHub Releases page](#) ²⁹ or install via Python's pip (not recommended unless you already have a Python setup) ³⁰. The makers of

`uv` suggest using `pipx` if you go the pip route (to keep it isolated) ³¹. For most beginners, the methods above are simpler.

2. Install `uv` on Windows: There are a few options for Windows: - **WinGet (Windows 10/11):** If you have the Windows Package Manager (WinGet) available, run PowerShell or Command Prompt as administrator and execute:

```
winget install --id=astral-sh.uv -e
```

This will download and install `uv` for you ³². - **Installer Script (Windows):** Alternatively, use PowerShell to run the official installer script:

```
PowerShell -ExecutionPolicy Bypass -Command "Invoke-RestMethod https://astral.sh/uv/install.ps1 | Invoke-Expression"
```

(You can shorten `Invoke-RestMethod` to the alias `irm` and `Invoke-Expression` to `iex` as in the docs ³³.) This will download the `uv` installer for Windows and execute it. You might need to approve running a remote script; the `-ExecutionPolicy Bypass` flag in the command handles that by allowing the script to run ³⁴. - **Scoop:** If you use [Scoop](#) on Windows, you can install `uv` by:

```
scoop install main/uv
```

³⁵. - **Manual:** As with Mac, you could download a release binary (`uv.exe`) from GitHub and put it in your PATH.

After installation, open a new terminal (or PowerShell) and verify it works by checking the version:

```
$ uv --version
uv 0.x.y (...commit or build info...)
```

If you see a version string, you're ready to go.

No Python? No problem! One of the most impressive features of `uv` is that it can bootstrap Python for you. You do **not** need to separately install Python from python.org or the Windows Store if you don't want to. For example, if you run:

```
$ uv venv
```

in a new directory, `uv` will detect that no Python is available and **automatically download the latest Python interpreter** to create a virtual environment ³⁶. By default, `uv` picks the newest stable Python version for new environments. You can also explicitly install specific Python versions via `uv` (e.g., `uv python install 3.11`), but it's optional. This means after installing `uv`, you can immediately start managing environments and running scripts without any additional setup – a huge win for beginners who might otherwise struggle with installing Python itself. *(If you prefer to manage Python manually, `uv` can use existing installations as well ³⁷, but the automatic approach is the default.)*

Shell Integration: `uv` includes shell tab-completion scripts for convenience. If you want, you can enable autocompletion so that commands and options tab-complete in Bash, Zsh, PowerShell, etc. The installation script may already add this, but if not, you can run a command like `uv generate-shell-completion bash` and add it to your shell's rc file ³⁸. This is optional but helpful as you learn `uv` commands.

Now that `uv` is set up (and you have Python through it), let's see it in action for managing dependencies in a simple scripting scenario.

Tutorial: Using uv for Standalone Python Scripts

In this section, we'll go through a beginner-friendly example of writing a Python script and using `uv` to handle its dependencies and execution. This is especially useful if you're leveraging AI assistance to generate code: often the AI will use libraries (like `requests`, `pandas`, etc.), and `uv` can automatically fetch those so your script runs without manual tinkering. The goal is to have a *single script that "just works"* when run with `uv`.

Writing and Running a Basic Script (No Dependencies)

First, let's start with the simplest case – a script that only uses Python's standard library. Create a file called `hello.py` with the following content:

```
# hello.py
print("Hello, world!")
```

If you run this with the normal Python interpreter, it would just print "Hello, world!". Let's run it with `uv`:

```
$ uv run hello.py
Hello, world!
```

That's it – `uv` noticed that your script had no special dependency requirements, so it just executed it in an isolated environment using the default Python ³⁹. Even though this script didn't need any extra packages, `uv` still ran it in a managed way (creating a temporary environment). If the script uses only standard library modules (or no imports at all), you don't need to do anything but `uv run`. You can also pass command-line arguments after the script name, and `uv` will forward them. For example:

```
$ uv run hello.py Alice Bob
Hello, world!
```

(assuming the script was modified to greet arguments from `sys.argv`). In short, `uv run <script.py>` is a general replacement for `python <script.py>`, with the difference that it sets up an environment for the script.

Note: If you run `uv run` inside a directory that is a `uv` project (i.e., contains a `pyproject.toml`), by default it will assume you might want to use the project's

environment ⁴⁰. In our case, we're working with standalone scripts (no project file), so this isn't an issue. But if you ever need to force `uv` to ignore a project and run the script independently, you can add the `--no-project` flag (placed before the script name) ⁴¹.

Adding Dependencies to a Script with Inline Metadata

Now, let's say we want to expand our script to do something more interesting that requires an external library. For example, maybe we want to fetch data from a web API and pretty-print some results. We decide to use the `requests` library for HTTP and the `rich` library for nice formatting.

A naive approach would be to just import those and try to run the script, but you'd get a `ModuleNotFoundError` because they're not installed:

```
# example.py
import time
from rich.progress import track
import requests

for _ in track(range(5), description="Fetching data..."):
    time.sleep(0.5)
response = requests.get("https://httpbin.org/ip")
print("Your IP is:", response.json().get("origin"))
```

If we try `uv run example.py` without specifying dependencies, it will fail, as expected:

```
$ uv run example.py
Traceback (most recent call last):
  File "/Users/me/example.py", line 3, in <module>
    from rich.progress import track
ModuleNotFoundError: No module named 'rich'
```

We have a couple of ways to tell `uv` what dependencies this script needs. One quick method is to specify them at run time with `--with`:

```
$ uv run --with rich --with requests example.py
```

This would instruct `uv` to ensure `rich` and `requests` are present before running. `uv` would resolve and install those packages into an isolated environment, then execute the script ⁴² ⁴³. However, typing out `--with ...` each time isn't ideal, and it doesn't "remember" the dependencies next time. A better approach is to **declare the requirements in the script itself**, so you only have to set it up once.

`uv` supports **inline metadata** for scripts, following a proposed standard (PEP 723) for embedding dependency information in Python files. You don't need to manually write this metadata; `uv` can add it for you. Let's use `uv` to add our needed packages to the script:

```
$ uv add --script example.py "requests<3" rich
Updated `example.py`
```

We specified `"requests<3"` and `"rich"` as the dependencies (meaning we want any 2.x version of requests, not a hypothetical future requests 3.x, just as an example of version specifier). After running that command, open **example.py** – you'll see a new annotated section at the top:

```
# /// script
# dependencies = [
#   "requests<3",
#   "rich",
# ]
# ///
import time
from rich.progress import track
import requests
...
```

`uv` added a special comment block that declares the script's dependencies in TOML format ⁴⁴. This block doesn't affect normal Python execution (since it's in comments), but `uv` knows to look for it. It's essentially equivalent to having a mini **pyproject** for this single file.

Now we can simply run the script without `--with`, as `uv` will read the inline metadata:

```
$ uv run example.py
Reading inline script metadata from: example.py
Resolved 5 packages in [few] ms
Installed 5 packages in [few] ms
Fetching data... 100% ██████████ 5/5 [00:02<00:00, 2.00it/s]
Your IP is: 203.0.113.45
```

On the first run, you'll notice `uv` reports installing some packages. We asked for `requests` and `rich`, but it says 5 packages were installed. That's because those libraries have their own dependencies (for example, `requests` typically brings in `urllib3`, `certifi`, etc.). `uv` resolved the whole dependency graph extremely quickly (a few milliseconds) and installed everything needed ⁴⁵ ⁴⁶. The output of our script (a progress bar from rich and then an IP address printout) is shown after the installation step.

If we run the script again, `uv` will see that the dependencies haven't changed and that it has a cached environment for this script, so it **won't reinstall** anything on subsequent runs. It will directly execute the script, making reuse very fast. In fact, `uv` caches environments keyed by the script content and requested dependencies. As long as the `# /// script` section and the chosen Python version remain the same, `uv` reuses the environment. This is how `uv` provides the convenience of ephemeral environments without the cost each time – *it's cached and only updated if you change the metadata*. If you do change the dependency list (say you add another library), `uv` will resolve the difference and install the new package on the next run.

A couple of things to note about inline metadata: - The special markers `# /// script` and `# ///` delineate the metadata block. `uv` expects a `dependencies` field there (even if it's an empty list) ⁴⁷ ⁴⁸. You can also specify a required Python version if your script needs a certain minimum version. For example, you might add `# requires-python = ">=3.10"` in that block to indicate the script needs Python 3.10+ ⁴⁷. `uv` will then ensure it uses a suitable Python interpreter (downloading one if necessary) to run the script ⁴⁹. - You can lock the dependencies for the script to exact versions using `uv lock`, which creates a `uv.lock` file. This is more advanced, but it's useful if you want to snapshot the working versions so that even if upstream releases change, your script continues to use the known-good versions ⁵⁰. By default, without locking, `uv` will always fetch the latest versions that satisfy your criteria (e.g., the newest `requests` < 3). For learning purposes, you might not worry about locking, but it's good to know for production use.

With our `example.py` set up, we can share this single file with others. If they have `uv` installed, they can simply run `uv run example.py` and it will automatically pull in the needed packages and run, exactly as it does for us. This makes sharing simple scripts extremely convenient – no more instructing someone to “install X, Y, Z then run the script”; the script is self-contained with its requirements.

Using a Shebang (Optional)

On Unix-like systems (macOS/Linux), you can even make your script directly executable by adding a special shebang line at the top. For `uv`, the shebang would be:

```
#!/usr/bin/env -S uv run --script
```

Place that as the first line, followed by the `# /// script` block and your code. For example:

```
#!/usr/bin/env -S uv run --script
# /// script
# dependencies = [
#     "requests<3",
#     "rich",
# ]
# ///
import time
...
```

After adding the shebang, give execution permission to the file (`chmod +x example.py` on Mac/Linux). Now you can run `./example.py` directly from the shell, and it will invoke `uv` under the hood to run your script ⁵¹ ⁵². This trick isn't needed for Windows (and Windows has a different approach to executable scripts), but it's a neat convenience on other systems for frequently used scripts.

Now that we've seen `uv` in action for plain `.py` files, let's explore how it fits into other common development setups like Jupyter Notebooks and VS Code.

Working with uv in Jupyter Notebooks

Jupyter Notebooks are popular for interactive Python, especially in data science. Typically, if you wanted to use a new library in a Jupyter notebook, you'd have to install it (either via command line or using

`%pip install` in a notebook cell) and manage which environment the Jupyter kernel is using. `uv` can simplify this in a couple of ways:

- **Quick ad-hoc Notebook:** If you just want to fire up a notebook server and play around (perhaps to test a snippet with certain libraries), you can use `uv` to launch Jupyter in an isolated environment on the fly. For example:

```
$ uv tool run jupyter lab
```

This command (equivalent to `uvx jupyter lab`) will ensure Jupyter is installed and then start JupyterLab ⁵³. It does *not* require you to have Jupyter pre-installed – `uv` will fetch the `jupyter` package if needed. The Jupyter server it launches will be running in a temporary environment separate from your system. You'll see the familiar Jupyter interface open (likely at `http://localhost:8888`). You can create notebooks and use them as usual. If you only need to use standard libraries or anything you explicitly installed in that session, you're all set. Once you shut down Jupyter, that environment can be discarded (or cached for the next `uv tool run` invocation of Jupyter).

- **Using `uv` in an existing environment or project:** Suppose you have a `uv` project or a persistent environment and you want to use Jupyter with it. In that case, you'd want Jupyter to run *with the same packages* you have in your environment. One approach is:
- Install Jupyter into your environment (so the kernel has access to all needed packages). If you have a `pyproject.toml` (project), you might do: `uv add --dev jupyter` (marking it as a development dependency). For a quick script environment, you could do: `uv pip install jupyter` inside the environment.
- Launch Jupyter via `uv run`. For example, within a project directory:

```
$ uv run --with jupyter jupyter lab
```

This ensures that the Jupyter server runs with the project's virtual environment loaded ⁵⁴. Then you can import your project's packages in the notebook. If your project had, say, `requests` as a dependency, doing `import requests` in a notebook cell will use the version from your project's `.venv`.

- If you need to add new packages while inside the notebook, you have options:
 - Use the notebook magic commands: e.g. `!pip install somepackage`. But be careful: if you launched Jupyter with `uv run --with jupyter`, the server is running in an isolated env separate from your project's env (unless you did an explicit integration). In such a case, `!pip install` in a cell might install into the Jupyter server's env, not the project. A more reliable way, if your notebook kernel is the project env, is to use `uv` inside the notebook. For example, you can run `!uv add newlib` in a notebook cell to add a new library to the project's dependencies (this updates `pyproject/uv.lock` and installs into the venv) ⁵⁵ ⁵⁶. Or `!uv pip install newlib` to just install it into the venv without adding to the project file ⁵⁷.
 - In short, if you've set up a **kernel** for the project's env, you can manage packages via `uv` from inside the notebook. Setting up a kernel is as easy as installing `ipykernel` in the env (`uv add --dev ipykernel`) and then running `uv run ipython kernel`

`install ...` to register the kernel ⁵⁸. In VS Code, you might not even need to do that manually if you just select the Python interpreter (more on that shortly).

For beginners, a simpler route is: use `uv` to start Jupyter for quick experiments, but do heavier work in scripts/projects. If you do prefer notebooks, you can still rely on `uv` to manage the environment that Jupyter is using. The key takeaway is that `uv` **can install Jupyter and any libraries you need on-the-fly**, so you don't have to fuss with pip/conda to get your notebook environment set up. For example, if you just learned about a new library `X` and want to try it in a notebook:

```
$ uv tool run jupyter notebook
```

Then in a notebook cell:

```
!uv pip install X
import X
```

This will install `X` into the running environment and you can use it immediately ⁵⁹. When you close the notebook server, that environment can be thrown away unless you saved it somewhere.

Using uv with VS Code (or Other Editors)

Visual Studio Code is a popular editor for Python, and many beginners use it along with its Python extension. The **Cursor** editor is another example, which integrates AI assistance for coding. Regardless of editor, the main question is how to use `uv` in your development workflow there. Here are some tips:

- **Running scripts via Terminal:** The simplest way is to use VS Code's integrated terminal to run `uv` commands. For instance, if you have `example.py` open, you can open a terminal (Ctrl+` in VS Code) and just type `uv run example.py`. This will execute the script as described earlier. This approach doesn't require any special editor configuration - you're basically using `uv` as you would in a normal shell. It's a good starting point for beginners: write code in the editor, run it with `uv` in the terminal, see results, repeat.
- **Using VS Code's Run/Debug buttons:** VS Code typically wants to know which Python interpreter to use for running and debugging. If you want to use `uv`'s managed environment, there are a couple of approaches:
- **Use a `uv`-created virtualenv as the interpreter:** You can create a persistent virtual environment for your project or script using `uv` and then point VS Code at it. For example, if you're working on a project (say you did `uv init myproj`), `uv` will likely create a `.venv` directory for that project with the environment. In our earlier script example, if you wanted a persistent env, you could do:

```
$ uv venv          # creates a .venv folder with the default Python
$ uv pip install requests rich # install needed packages into .venv
```

(The second step is analogous to what `uv add --script` did, but this time we're manually installing into a persistent env.) In VS Code, you can then use the command "Python: Select Interpreter" and choose the Python from your project's `.venv` (it should appear in the list if VS Code detects it, or you can navigate to `.venv/bin/python` on Mac/Linux or `.venv\Scripts\python.exe` on Windows) ⁶⁰ ⁶¹. Once selected, the VS Code run button (▶) or debugger will use that interpreter, which has your dependencies. Essentially, you're using `uv` to set up the env, then using that env in VS Code like any normal virtualenv. The advantage is `uv` ensured you got the right packages and Python version quickly.

- **Configure VS Code to use `uv` directly:** This is a bit advanced and not usually necessary, but VS Code's launch configurations could be set to call `uv run` instead of `python`. For instance, you could make a debug launch config that runs the command `uv` with args `run yourscript.py`. However, debugging might not attach easily in that scenario since `uv` spawns its own process. A simpler path is the previous one: use a `uv` environment as the interpreter so that VS Code's own Python extension can work with it.
- **Editor Intellisense and Imports:** When you use `uv` inline metadata for a script, VS Code (or other editors) might not know about the dependencies for autocompletion, because it doesn't read those comment lines. If you want IDE features like IntelliSense to recognize imported packages, you should have the packages installed in an environment that the editor knows about. One workaround is to manually create a venv as above and install the packages, or use `uv` to do so. If you've run `uv run example.py` once, the packages are installed *somewhere*, but VS Code won't automatically use that ephemeral env. In such cases, you might decide to convert your script into a small project:
 - Run `uv init .` in the folder to create a `pyproject.toml` for it, add dependencies with `uv add`, which will make a `.venv` and install them. Then VS Code will detect that `.venv`.
 - Alternatively, you can use the **"Python: Create Environment"** feature in VS Code to create a venv and then use `uv pip install` to install packages into it. For example, if VS Code makes a `.venv`, you could open a terminal and run `uv pip install requests rich` while in that environment (ensuring the `UV_HOME` or so is pointed to that, but if you activated the venv, `uv pip` should default to it because `UV_PYTHON` might pick the active interpreter – this is a bit internal detail).

In summary, **for a beginner** it's perfectly fine to manage packages with `uv` but use the editor's Python extension to run the code via a venv. `uv` is compatible with that because it produces standard virtual environments. The key is to let `uv` handle installing Python and packages, then point your editor to the right interpreter. Many beginners find it easiest to just use the terminal though, and that's okay too.

Using Cursor: If you're using Cursor or another AI coding editor, the workflow is similar. You can ask the AI to write a script, then run it with `uv`. If the AI writes `import something` that isn't built-in, you don't manually fiddle with pip – just run with `uv` and it will grab it. This tight loop allows you to focus on coding (possibly with AI help) rather than on dependency logistics.

Example VS Code Workflow (summary)

1. **Create a project/env:** Let's say you want to build a small project in VS Code. You could do:

```
$ uv init myproject
$ cd myproject
$ uv add numpy pandas # for example, add some deps
```

This will create a directory `myproject` with a `pyproject.toml` listing `numpy` and `pandas`, and a fresh `.venv` containing those packages ⁶². Open this folder in VS Code (`code .` from that directory).

2. **Select interpreter:** In VS Code, select the interpreter from `myproject/.venv`. Now when you create a new Python file or Jupyter notebook in VS Code, it will use that env (you should see “Python 3.x (.venv)” in the status bar).
3. **Write and run code:** Now you can import `numpy`, `pandas`, etc., and VS Code will have IntelliSense for them because it knows about the env. Running the code (via run button or F5 for debug) will execute in that env – which was set up by `uv`. If you need a new package, you can either:
 4. Use VS Code’s integrated terminal: `uv add package` (to add to `pyproject` and install) or `uv pip install package` (to just install into the env without tracking).
 5. Or use the notebook magic if in a `.ipynb`: `!uv add package` in a cell.
 6. The environment updates, and VS Code should pick up the changes (you might sometimes need to reload the window for IntelliSense, but generally it should work).

This approach combines `uv`’s power with VS Code’s convenience.

Traditional Tools Recap (pip, pipx, Poetry, Conda) vs uv

Before we conclude, it’s worth summarizing how `uv` stands relative to the traditional tools we discussed, as this cements why one might choose `uv` as a beginner:

- **pip + venv:** Manual but standard. You must manage environments and dependencies yourself. No explicit locking unless you add other tools. `uv` automates the environment creation and uses a faster installer, while still being compatible with `pip`’s conventions (it can read `requirements.txt` or output one). If `pip/venv` is like a basic car, `uv` is a high-speed automated vehicle – you give it the route (requirements) and it takes you there faster and with less hassle.
- **pipx:** Great for installing runnable tools globally. `uv`’s `tool` commands cover this use case, and actually `uv` can manage those global tools with the same speed. For instance, installing a tool with `uv tool install` is nearly instant ⁶³ ⁶⁴. If you were, say, instructed to install `black` (a code formatter) or `httpie` (a CLI HTTP client), you could use `pipx` or simply do `uv tool install black`. Both achieve isolation, but `uv` keeps everything within one system (and one cache).
- **Poetry:** A one-stop solution for project dependency management with `pyproject.toml`. `uv` offers the same in-project experience (with `uv init`, `uv add`, `uv lock`, etc.) but is generally much faster and has extra features (like easily managing multiple Python versions). One article noted that *while Poetry was a significant improvement over traditional tools, UV takes performance to another level* ⁶⁵. Also, `uv`’s design to be compatible with `pip` tools offers flexibility that Poetry’s more rigid approach doesn’t ⁶⁵. For a beginner, Poetry is friendly but might feel slow or heavy at times; `uv` might feel snappier and more integrated with the system.

- **Conda:** Solves some issues that pip doesn't (like binary dependencies, system libraries, etc.). `uv` is focused on Python packages (it doesn't install, say, an OS-level library or GPU driver – conda would, for example, handle installing CUDA toolkit along with TensorFlow from its channels). If you are in a domain where you need those capabilities, conda or mamba might still be needed. But for a huge range of beginner projects (web, scripts, data analysis with pure Python libs), `uv` is sufficient and avoids the bulk of installing Anaconda. `uv` also works with PyPI by default, which has the most Python packages. In terms of environment management, both `uv` and conda handle creating and switching envs; conda is cross-language, but `uv` is laser-focused on Python. Notably, conda environments aren't as **lightweight** as `uv`'s (conda itself has a heavier footprint ⁶⁶). `uv` also yields consistency across platforms (the same lockfile can be used on Windows, Mac, Linux to install from PyPI) whereas pip can sometimes have platform-specific wheels – but `uv`'s resolver can produce “universal” requirements files that work on any OS ⁶⁷.

To sum up, `uv` **doesn't necessarily replace every scenario** (for instance, conda's niche in system packages), but it covers the vast majority of Python dependency management needs in one tool. It gives beginners a clear path: *install uv, and use it for everything from installing Python itself to adding packages, running scripts, and managing project environments*. This means less context-switching between tools and fewer chances to make mistakes.

Conclusion and Next Steps

Dependency management in Python has a reputation for being difficult, but with modern tools like `uv`, it's becoming much more approachable. In this guide, we discussed how traditional methods require juggling multiple utilities and concepts (which can be daunting for newcomers) and how `uv` streamlines these tasks. We walked through setting up `uv` on both Mac and Windows, so you can get started even on a fresh machine without Python. We then demonstrated using `uv` to run a simple Python script, gradually adding dependencies and showing how `uv` handles installation and caching automatically. We also touched on using `uv` in interactive notebooks and IDEs, which is important for a well-rounded learning experience.

By using `uv` as your primary package manager, you gain: - **Ease of use:** one command to run scripts with all dependencies handled. - **Speed:** super-fast installs and updates, so you spend less time waiting and more time coding. - **Reproducibility:** the option to lock down versions and the confidence that `uv` will create isolated environments, ensuring that your script works “out of the box” for others. - **Less frustration:** no more “it works on my computer but not on theirs” or puzzling over whether you installed a library in the correct environment – `uv` manages those details for you.

As a beginner focusing on Python scripting (especially with the assistance of AI to generate code), adopting `uv` can smooth out the learning curve. Instead of wrestling with environment issues, you can concentrate on the fun part: writing and running Python code. And if you ever need to dive deeper, `uv` has extensive documentation on advanced features like workspace management, publishing packages, and integration with CI/CD ²⁰ ¹⁰ – capabilities that you can gradually explore as you become more confident.

Next Steps: Try converting one of your existing small projects or scripts to use `uv`. For instance, take a script that you previously ran with system Python and pip, and run it with `uv` by adding the necessary `uv` metadata. Observe how `uv` creates an environment and installs packages quickly. You can also experiment with `uv`'s other commands: use `uv python install` to grab a new Python version, or `uv pip freeze` to see what's in the environment. With a bit of practice, you'll develop an intuitive sense for how `uv` simplifies dependency management.

By leveraging `uv`, you have a potent tool in your arsenal that will serve you well as you progress from beginner to proficient Pythonista. Happy scripting!

Sources:

- Python dependency management challenges and uv's goals 5 11
- uv's speed and unified features 7 19
- uv automatic virtual env handling and isolation 10 1
- Installing and using uv on macOS and Windows 26 32
- uv can auto-download Python if not present 15
- Example of adding inline script dependencies and running with uv 44 46
- Using uv with Jupyter and VS Code environments 68 69
- Comparison of uv with pip/conda/poetry (features and trade-offs) 2 3

1 39 40 41 42 43 44 46 47 48 49 50 51 52 Running scripts | uv
<https://docs.astral.sh/uv/guides/scripts/>

2 3 4 5 11 65 66 Python UV: The Ultimate Guide to the Fastest Python Package Manager | DataCamp
<https://www.datacamp.com/tutorial/python-uv>

6 7 8 9 12 13 14 16 17 18 45 62 63 64 67 GitHub - astral-sh/uv: An extremely fast Python package and project manager, written in Rust.
<https://github.com/astral-sh/uv>

10 19 20 21 Managing Python Projects With uv: An All-in-One Solution – Real Python
<https://realpython.com/python-uv/>

15 36 37 Installing and managing Python | uv
<https://docs.astral.sh/uv/guides/install-python/>

22 23 24 Caching | uv
<https://docs.astral.sh/uv/concepts/cache/>

25 26 27 28 29 30 31 32 33 34 35 38 Installation | uv
<https://docs.astral.sh/uv/getting-started/installation/>

53 54 55 56 57 58 59 60 61 68 69 Using uv with Jupyter | uv
<https://docs.astral.sh/uv/guides/integration/jupyter/>